

Spacez — AI Product Associate · Take-Home

Time: One day. Your submission is due **24 hours after you receive this brief**. We care how you think, not how polished it looks.

Context

Spacez runs premium villa stays. Bookings come via Airbnb, Booking.com, and Google, and guests review on whichever platform they booked through. Three groups care about those reviews:

- **Operations** — recurring fixable issues (maintenance, housekeeping, check-in).
- **Business** — portfolio performance, pricing, reputation, which properties to back or drop.
- **Caretakers** — the on-ground hosts and the guest's main point of contact.

[**VERIFY — replace with Spacez's real caretaker model**] Caretakers are assigned to a property (some cover more than one nearby), handle check-in/checkout and guest requests, and coordinate housekeeping/cooks (often third-party vendors).

The problem

Reviews are scattered across three platforms and nobody synthesizes them — so issues recur, hospitality quality is invisible at scale, and the business can't see patterns. We want an **AI agent that reads reviews from all platforms and turns them into something each group can act on**.

Your task

1. **Design the full concept** for all three stakeholders: what the agent outputs, who acts, what action it drives.
2. **Build a working prototype for ONE stakeholder** (your pick — tell us why), deeply. Reason about the other two in the spec.

Pressure-test this, don't assume it: we currently think caretakers should get reports on reviews from bookings they hosted. Good idea or risky? Show your reasoning.

Data

Use the **attached Excel file**: ~46 synthetic-but-realistic reviews across Airbnb, Booking.com, and Google. Sentiment, themes, and normalisation are deliberately *not* pre-computed — that's the work. Extend it if useful (label additions).

Deliverables

1. **Prototype** — a link we can open (any tool).
2. **Spec (max 2 pages)**: your problem framing; the flow (trigger → analysis → output → who acts → action); what each stakeholder gets and why; key trade-offs; how you'd measure

success.

3. **Risks & pushback (½ page):** where this is weak and what you'd validate first.
4. **Walkthrough (5-10 min Loom):** your process, prompts, what you dropped, where the AI was wrong.

Constraints

- Any AI build tool. **Free tier is enough; tooling and spend won't be scored.**
- Use the attached dataset.

What we're looking for

Right-problem judgment · outputs that drive *different* actions per stakeholder · how you used AI (not just the output) · an executable spec · willingness to challenge the brief. **A scrappy prototype with sharp thinking beats a polished one with shallow thinking.**

Submission

Share a single **Google Doc** (with view access) containing your spec, your risks & pushback section, the prototype link, and the Loom link. **Due 24 hours after you receive this brief.**

INTERNAL — REMOVE BEFORE SENDING

Score 1-5 each: problem framing · stakeholder reasoning (distinct actions, not three identical reports) · handling the caretaker hypothesis (separates caretaker-controllable from property/listing issues) · data judgment · AI process · communication · pushback.

Did they catch the seeded traps?

- Rating scales differ (Booking /10, Airbnb & Google /5) — normalise, don't blind-average.
- Many low ratings aren't the caretaker's fault (wifi, oversold photos, bad road, gate policy).
- Lokesh Gowda covers two properties — check-in delays follow *him*, not the property.
- Recurring issues (Serenity pool, Cliffside heating, Vineyard cleanliness) vs. one-offs.
- Noise ("Amazing.") and the occupancy-policy review ≠ hospitality failures.

Quick call: caught ~0-1 traps + accepted the caretaker premise uncritically → likely no. Caught most + challenged the premise well → strong, even if the prototype is rough.